

vowel_duration_python

May 12, 2022

0.1 Materials

At the start, we have audio and annotated textgrids of **regilaul** songs, annotated for ictus/off-ictus and phrase text, then force-aligned using Praat's built in eSpeak forced aligner for Estonian to word and then segment. Then, we use the `estnltk.vabamorf` package to syllabify the words so that we can annotate the textgrid further with syllable quantity (Estonian has 3) and whether or not it is accented at the word level. We end up with a dataframe containing the data from three of the (Interval) tiers of the textgrid, acquiring duration data for words, individual segments, and (eventually) syllables.

```
[ ]: import parselmouth
from estnltk.vabamorf.morf import syllabify_word
import tgt
import string
#test method on a single TextGrid:
gridDir2 = "/Users/sarah/qp_final/txtgridtest/69.TextGrid"

def get_duration_labels(textgrid, tiername1, tiername2, tiername3):
    tmp = tgt.read_textgrid(textgrid)
    mytier = tmp.get_tier_by_name(tiername1)
    other = tmp.get_tier_by_name(tiername2)
    ictus = tmp.get_tier_by_name(tiername3)
    segments = []
    word_dur = mytier.intervals

    for interval in word_dur:
        h = interval.start_time
        t = interval.end_time
        b = t-h
        l = interval.text
        tmpseg = [other.get_annotations_between_timepoints(h,t)]
        s = syllabify_word(l, as_dict=True)
        i = 0
        for list in tmpseg:
            n = 0
            while i < len(s):
                item = s[i]
```

```

        shape = item.get('syllable')
        q = item.get('quantity')
        a = item.get('accent')
        geminate = False
        for char in shape.strip(string.punctuation):
            if geminate: continue
            if n < len(list):
                myinterval = list[n]
                c = myinterval.text
                d = myinterval.start_time
                g = myinterval.end_time
                tmpick = ictus.
↪get_annotations_between_timepoints(d,g,left_overlap=True,right_overlap=True)
                if len(tmpick) == 0: ick = "off"
                else: ick = "ictus"
                #segment duration
                j = g-d
                #segment midpoint for later measurements
                mid = g - (j/2)
                if " " in c:
                    row = [l,b,shape,q,a,c,j,mid,ick]
                    segments.append(row)
                    geminate = True

                elif char == c:

                    row = [l,b,shape,q,a,c,j,mid,ick]
                    segments.append(row)
                else :
                    can = "(" + char + ")" + c
                    row = [l,b,shape,q,a,can,j,mid,ick]
                    segments.append(row)
            n+=1
        i+= 1

    nu_df = pd.
↪DataFrame(segments,columns=["word","word_dur","syll","quantity","stress","segment","seg_dur"]
    return nu_df

syl_dur_df = get_duration_labels(gridDir2,"word","word/phon","ictus")
syl_dur_df.head()

```

```
[ ]:      word  word_dur syll  quantity  stress segment  seg_duration  \
0      Oh  0.240882  oh          3        1      o    0.165043
1      Oh  0.240882  oh          3        1      h    0.075839
2      me  0.169500  me          3        1      m    0.059664
3      me  0.169500  me          3        1      e    0.109836
4  vaesed  0.619500  vae          2        1      v    0.101258

      seg_midpoint  ictus
0      0.082521  ictus
1      0.202962  ictus
2      0.270713  ictus
3      0.355463  ictus
4      0.461010  ictus
```

0.2 Adding Spectral data

now that we have the duration data from the textgrid, we can query specific timepoints for information about the acoustic signal. The following function uses the midpoint (which we snagged while we were making the dataframe above) and get the first three formants(Hz) for each segment.

```
[ ]: import parselmouth

test = "/Users/sarah/qp_final/wavs/069.wav"

def get_formants(syl_dur_df, wave):
    song = parselmouth.Sound(wave)
    formant = song.to_formant_burg()
    f1 = []
    f2 = []
    f3 = []
    for float in syl_dur_df.seg_midpoint:
        time = float
        f1.append(formant.get_value_at_time(1,time))
        f2.append(formant.get_value_at_time(2, time))
        f3.append(formant.get_value_at_time(3,time))
    syl_dur_df["f1"] = f1
    syl_dur_df["f2"] = f2
    syl_dur_df["f3"] = f3
    return syl_dur_df
nu_df = get_formants(syl_dur_df,test)
nu_df.head()
```

```
[ ]:      word  word_dur syll  quantity  stress segment  seg_duration  \
0      Oh  0.240882  oh          3        1      o    0.165043
1      Oh  0.240882  oh          3        1      h    0.075839
2      me  0.169500  me          3        1      m    0.059664
3      me  0.169500  me          3        1      e    0.109836
4  vaesed  0.619500  vae          2        1      v    0.101258
```

	seg_midpoint	ictus	f1	f2	f3
0	0.082521	ictus	540.905421	989.905678	1414.238432
1	0.202962	ictus	850.865128	1620.784985	2186.523346
2	0.270713	ictus	346.158399	1818.412065	2798.007639
3	0.355463	ictus	411.766557	1385.397329	2214.984918
4	0.461010	ictus	488.221699	1013.678738	1576.019602

```
[ ]: from os.path import join
      #runs a for loop over a directory using the above-specified functions

test = "/Users/sarah/qp_final/txtgridtest/"
songs = "/Users/sarah/qp_final/songs/"

for fn in os.listdir(test):
    if '.TextGrid' not in fn:
        continue
    n = fn.strip('.TextGrid')
    wave = join(songs, n + '.wav')
    data_file = open("ictus_forms_" + n + ".csv", 'w')
    #make a dataframe with the interval tiers of the textgrid
    tmp = pd.DataFrame(get_duration_labels(join(test,fn), "word", "word/
↳phon", "ictus"))
    #add the formant data to the dataframe
    nu_df = get_formants(tmp, wave)
    print(nu_df.head())
    # nu_df.to_csv(data_file)
    # data_file.close()
```

	word	word_dur	syll	quantity	stress	segment	seg_duration	seg_midpoint	\
0	Löpe,	0.877833	lō	2	1	l	0.102860	0.206697	
1	Löpe,	0.877833	lō	2	1	(ō)	0.293082	0.404668	
2	Löpe,	0.877833	pe,	2	0	p	0.191140	0.646779	
3	Löpe,	0.877833	pe,	2	0	e	0.290751	0.887725	
4	lōpe,	0.836977	lō	2	1	l	0.075601	1.070901	

	ictus	f1	f2	f3
0	ictus	335.422530	1063.568157	1890.268233
1	ictus	684.001970	1113.986905	1962.573748
2	ictus	408.802542	1208.989186	2321.386963
3	off	758.503513	1717.665908	2219.432604
4	ictus	772.408953	1868.617726	1966.346496

	word	word_dur	syll	quantity	stress	segment	seg_duration	seg_midpoint	\
0	miks	0.334037	miks	3	1	m	0.109658	2.583638	
1	miks	0.334037	miks	3	1	i	0.103965	2.690450	
2	miks	0.334037	miks	3	1	k	0.071640	2.778252	
3	miks	0.334037	miks	3	1	s	0.048775	2.838459	
4	sa	0.174934	sa	3	1	s	0.049593	2.887643	

	ictus	f1	f2	f3
0	ictus	1015.032176	1668.628090	2451.566636
1	ictus	1018.542344	1242.393701	2614.593523
2	ictus	739.889853	1206.829336	2557.715244
3	ictus	1219.241465	1330.686940	2551.563447
4	off	996.088003	1411.057879	2533.755085

	word	word_dur	syll	quantity	stress	segment	seg_duration	seg_midpoint	\
0	Kelle	0.54837	kel	2	1	k	0.004193	0.002097	
1	Kelle	0.54837	kel	2	1	e	0.224554	0.116470	
2	Kelle	0.54837	kel	2	1	l	0.026000	0.241747	
3	Kelle	0.54837	le	1	0	l	0.105220	0.307357	
4	Kelle	0.54837	le	1	0	e	0.188403	0.454169	

	ictus	f1	f2	f3
0	ictus	NaN	NaN	NaN
1	ictus	1110.553772	2154.184673	3278.566300
2	off	723.585094	1951.030424	3121.543094
3	off	956.257232	1943.928928	3170.558668
4	off	699.571764	1442.320280	2391.975132

	word	word_dur	syll	quantity	stress	segment	seg_duration	\
0	Kuus,	0.690416	kuus,	3	1	k	0.142235	
1	Kuus,	0.690416	kuus,	3	1	u	0.232104	
2	Kuus,	0.690416	kuus,	3	1	u	0.161435	
3	kuus,	0.661647	kuus,	3	1	k	0.100761	
4	kuus,	0.661647	kuus,	3	1	u	0.293682	

	seg_midpoint	ictus	f1	f2	f3
0	0.071117	ictus	1028.357905	2289.070707	3278.949044
1	0.258286	ictus	1218.143866	2223.410006	3181.682865
2	0.455056	off	355.504806	1076.764366	1934.314752
3	0.740797	ictus	680.358222	1148.112933	1979.469484
4	0.938018	ictus	669.587600	1146.386219	1664.458201

	word	word_dur	syll	quantity	stress	segment	seg_duration	\
0	Laula,	0.616537	lau	2	1	l	0.138245	
1	Laula,	0.616537	lau	2	1	a	0.156972	
2	Laula,	0.616537	lau	2	1	u	0.077000	
3	Laula,	0.616537	la,	2	0	l	0.082000	
4	Laula,	0.616537	la,	2	0	a	0.162320	

	seg_midpoint	ictus	f1	f2	f3
0	0.441088	ictus	580.289721	1408.585781	2501.362541
1	0.588696	ictus	599.037686	1153.155750	1944.133038
2	0.705683	ictus	641.731079	1076.715564	2487.946642
3	0.785183	ictus	529.871527	1631.906656	1976.346982
4	0.907342	off	618.232042	1140.675032	1746.136521

	word	word_dur	syll	quantity	stress	segment	seg_duration	seg_midpoint	\
0	Laske	0.700087	las	2	1	l	0.068795	0.034398	

1	Laske	0.700087	las	2	1	a	0.221538	0.179564
2	Laske	0.700087	las	2	1	s	0.074939	0.327802
3	Laske	0.700087	ke	1	0	k	0.102666	0.416605
4	Laske	0.700087	ke	1	0	e	0.232149	0.584012

	ictus		f1		f2		f3	
0	ictus	1472.974201	2404.084622	3447.483334				
1	ictus	1125.029770	2441.781428	3331.398284				
2	ictus	854.440862	1641.705923	2960.910199				
3	ictus	686.227370	1665.848474	3056.221222				
4	ictus	895.423885	1680.387334	1883.316265				

	word	word_dur	syll	quantity	stress	segment	seg_duration	seg_midpoint	\
0	pandi	0.45674	pan	2	1	p	0.047328	6.835351	
1	pandi	0.45674	pan	2	1	a	0.177026	6.947528	
2	pandi	0.45674	pan	2	1	(n)	0.041698	7.056890	
3	pandi	0.45674	di	1	0	d	0.059916	7.107697	
4	pandi	0.45674	di	1	0	i	0.129039	7.202174	

	ictus		f1		f2		f3	
0	ictus	605.862912	1068.944908	2291.317753				
1	ictus	603.056801	1303.781807	2625.465878				
2	ictus	765.874507	1279.851220	2546.395265				
3	off	930.416284	1876.949848	2542.774896				
4	off	1074.849998	2167.095173	2671.989524				

	word	word_dur	syll	quantity	stress	segment	seg_duration	\
0	Oh	0.240882	oh	3	1	o	0.165043	
1	Oh	0.240882	oh	3	1	h	0.075839	
2	me	0.169500	me	3	1	m	0.059664	
3	me	0.169500	me	3	1	e	0.109836	
4	vaesed	0.619500	vae	2	1	v	0.101258	

	seg_midpoint	ictus		f1		f2		f3
0	0.082521	ictus	889.425959	2009.675238	3152.670429			
1	0.202962	ictus	639.966111	1118.912389	2929.645146			
2	0.270713	ictus	464.249379	840.111759	1332.896822			
3	0.355463	ictus	360.465479	1015.163348	1450.565657			
4	0.461010	ictus	504.486236	1459.123023	2328.845279			

	word	word_dur	syll	quantity	stress	segment	seg_duration	\
0	Ui-sui-sui-sui	1.297146	ui	3	1	u	0.103226	
1	Ui-sui-sui-sui	1.297146	ui	3	1	i	0.136909	
2	Ui-sui-sui-sui	1.297146	sui	3	1	s	0.065000	
3	Ui-sui-sui-sui	1.297146	sui	3	1	u	0.101500	
4	Ui-sui-sui-sui	1.297146	sui	3	1	i	0.138257	

	seg_midpoint	ictus		f1		f2		f3
0	0.051613	ictus	735.327729	2002.710633	2764.319990			
1	0.171681	ictus	934.145628	1833.703843	2616.755205			
2	0.272635	ictus	981.588844	2056.887357	2632.960528			

```

3      0.355885  ictus  958.855987  1841.503699  2488.421439
4      0.475764   off  470.683808  1001.828609  2784.254879

```

1 Now we put it into a big pile!

Here we concatenate all the data we have so far into one large pandas dataframe. At this point, we can keep annotating songs for the corpus, and as textgrids are finished we can run the scripts above to add them into the larger dataset. We're also gonna take the opportunity to add some metadata to the dataframes: fileid(song) and performer initials as potential grouping factors.

```

[ ]: import os
import pandas as pd
import statsmodels.formula.api as smf
folder = "/Users/sarah/qp_final/data_glob/"
meta = pd.read_csv("/Users/sarah/qp_final/song_metadata.csv")

songs_dfs = []
for fn in os.listdir(folder):
    if '.csv' not in fn: continue
    whole_name = os.path.join(folder,fn)
    song_df = pd.read_csv(whole_name)
    fileid1 = fn.strip('ictus_forms_')
    fileid = int(fileid1.strip('.csv'))
    row = meta.index[meta['track'] == fileid].tolist()
    performer = meta.performer[row[0]]
    for index in song_df:
        song_df['fileid'] = fileid
        song_df['performer'] = performer

    songs_dfs.append(song_df)

big_frame = pd.concat(songs_dfs, ignore_index=True)
print(big_frame.describe())
big_frame

```

	Unnamed: 0	word_dur	quantity	stress	seg_duration \
count	1967.000000	1967.000000	1967.000000	1967.000000	1967.000000
mean	198.710727	0.927221	1.887646	0.522623	0.139460
std	192.058558	0.467994	0.649035	0.499615	0.096460
min	0.000000	0.083718	1.000000	0.000000	0.004193
25%	56.000000	0.566898	1.000000	0.000000	0.076177
50%	131.000000	0.819799	2.000000	1.000000	0.109868
75%	259.000000	1.304484	2.000000	1.000000	0.180858
max	734.000000	2.168336	3.000000	1.000000	1.201817

	seg_midpoint	f1	f2	f3	fileid
--	--------------	----	----	----	--------

count	1967.000000	1966.000000	1966.000000	1966.000000	1967.000000
mean	57.490687	705.838851	1570.350446	2466.377994	43.916116
std	60.299049	202.075224	389.769265	331.870332	29.497152
min	0.002097	275.170888	519.638388	1204.360788	9.000000
25%	14.617182	548.215226	1296.586430	2290.649062	18.000000
50%	32.759004	694.356670	1529.016068	2496.860314	41.000000
75%	82.097020	855.569626	1831.690906	2669.113021	65.000000
max	219.349440	1472.974201	2693.060094	3552.062191	94.000000

```
[ ]:      Unnamed: 0      word  word_dur  syll  quantity  stress  segment  \
0          0      pandi  0.456740  pan          2          1          p
1          1      pandi  0.456740  pan          2          1          a
2          2      pandi  0.456740  pan          2          1      (n)
3          3      pandi  0.456740  di          1          0          d
4          4      pandi  0.456740  di          1          0          i
...      ...      ...      ...      ...      ...      ...
1962      34  sulased,  0.772673  la          1          0          l
1963      35  sulased,  0.772673  la          1          0      (a)
1964      36  sulased,  0.772673  sed,         3          1          s
1965      37  sulased,  0.772673  sed,         3          1          e
1966      38  sulased,  0.772673  sed,         3          1          d
```

	seg_duration	seg_midpoint	ictus	f1	f2	\
0	0.047328	6.835351	ictus	605.862912	1068.944908	
1	0.177026	6.947528	ictus	603.056801	1303.781807	
2	0.041698	7.056890	ictus	765.874507	1279.851220	
3	0.059916	7.107697	off	930.416284	1876.949848	
4	0.129039	7.202174	off	1074.849998	2167.095173	
...	
1962	0.056635	8.946193	ictus	935.586178	2102.100627	
1963	0.152365	9.050693	ictus	564.758290	1384.443183	
1964	0.093000	9.173375	ictus	1298.894426	1954.360862	
1965	0.144135	9.291943	off	909.702594	1747.600824	
1966	0.078979	9.403500	ictus	840.561088	1461.144276	

	f3	fileid	performer
0	2291.317753	65	LK
1	2625.465878	65	LK
2	2546.395265	65	LK
3	2542.774896	65	LK
4	2671.989524	65	LK
...
1962	3341.040339	69	L0
1963	3073.026293	69	L0
1964	3033.911307	69	L0
1965	2806.347885	69	L0
1966	1987.093716	69	L0

[1967 rows x 15 columns]

```
[ ]: #for the present paper, we are only interested in the vowel durations:
#set of Estonian vowels to filter the dataframe:
vowels = ["i"," ", "y","e"," ", "ø","æ","a"," ", "o"," ", "u"," " ]
vowel_df = big_frame[big_frame.segment.isin(vowels)].copy()
print(vowel_df.describe())
vowel_df.head()
```

	Unnamed: 0	word_dur	quantity	stress	seg_duration \
count	615.000000	615.000000	615.000000	615.000000	615.000000
mean	197.894309	0.949184	1.852033	0.528455	0.197959
std	190.819594	0.492294	0.684323	0.499596	0.117456
min	0.000000	0.083718	1.000000	0.000000	0.015260
25%	57.500000	0.555927	1.000000	0.000000	0.131685
50%	132.000000	0.846330	2.000000	1.000000	0.178758
75%	259.000000	1.335653	2.000000	1.000000	0.236604
max	734.000000	2.168336	3.000000	1.000000	1.201817

	seg_midpoint	f1	f2	f3	fileid
count	615.000000	615.000000	615.000000	615.000000	615.000000
mean	57.435001	720.172015	1565.957831	2467.122054	44.450407
std	59.849928	198.120808	396.035823	332.652632	29.328873
min	0.051613	275.499384	519.638388	1396.638475	9.000000
25%	15.556456	567.050872	1262.772049	2300.155478	18.000000
50%	33.763837	703.407720	1523.729941	2496.694808	41.000000
75%	78.970333	867.648661	1838.565263	2663.815935	65.000000
max	219.349440	1219.144093	2693.060094	3331.398284	94.000000

```
[ ]: Unnamed: 0    word word_dur syll quantity stress segment \
1          1    pandi  0.456740  pan         2         1        a
4          4    pandi  0.456740   di         1         0        i
6          6  mind(e)  0.292967  mind         2         1        i
9          9  mind(e)  0.292967   (e)         2         0        e
11         11    paju  0.881775   pa         1         1        a
```

	seg_duration	seg_midpoint	ictus	f1	f2	f3 \
1	0.177026	6.947528	ictus	603.056801	1303.781807	2625.465878
4	0.129039	7.202174	off	1074.849998	2167.095173	2671.989524
6	0.085210	7.364037	off	731.376613	1651.342030	2557.796933
9	0.067737	7.527525	off	852.719874	1444.954754	2241.764747
11	0.386617	7.887546	ictus	794.798321	850.845027	2525.457521

	fileid	performer
1	65	LK
4	65	LK
6	65	LK

```

9      65      LK
11     65      LK

```

```

[ ]: #make sure quantity is read as a categorical variable
vowel_df['quantity'] = vowel_df['quantity'].astype('object')
vowel_df['stress'] = vowel_df['stress'].astype('object')
print("vowel duration and stressed/unstressed: \n" , vowel_df.
      ↳groupby('stress')['seg_duration'].mean())
print("vowel duration and syllable quantity: \n" , vowel_df.
      ↳groupby('quantity')['seg_duration'].mean())
print("vowel duration and ictus/off-ictus \n" , vowel_df.
      ↳groupby('ictus')['seg_duration'].mean())

```

vowel duration and stressed/unstressed:

```

stress
0    0.228611
1    0.170609

```

Name: seg_duration, dtype: float64

vowel duration and syllable quantity:

```

quantity
1    0.216361
2    0.205660
3    0.140581

```

Name: seg_duration, dtype: float64

vowel duration and ictus/off-ictus

```

ictus
ictus    0.202398
off      0.191118

```

Name: seg_duration, dtype: float64

```

[ ]: import pandas as pd
import statsmodels .formula.api as smf

#is ictus (song prominence) a good predictor for vowel duration?

ickmodel = smf.ols("seg_duration ~ ictus", data=vowel_df).fit()
ickmodel.summary()

```

```

[ ]: <class 'statsmodels.iolib.summary.Summary'>
"""

```

```

                                OLS Regression Results
=====
Dep. Variable:                seg_duration    R-squared:                0.002
Model:                        OLS            Adj. R-squared:         0.001
Method:                      Least Squares   F-statistic:            1.355
Date:                        Thu, 12 May 2022  Prob (F-statistic):    0.245
Time:                        00:26:19        Log-Likelihood:         445.67

```

No. Observations: 615 AIC: -887.3
 Df Residuals: 613 BIC: -878.5
 Df Model: 1
 Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.2024	0.006	33.290	0.000	0.190	0.214
ictus[T.off]	-0.0113	0.010	-1.164	0.245	-0.030	0.008
Omnibus:	496.206	Durbin-Watson:	1.603			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	11531.163			
Skew:	3.442	Prob(JB):	0.00			
Kurtosis:	23.065	Cond. No.	2.44			


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 ""

well, the intercept coefficient is significant ($p < 0.05$), so vowels that are in the ictus position are predictably longer than those in the off-ictus or weaker positions in the song. The R squared is still pretty small, though. Let's see the residuals.

```
[ ]: ickmodel.fittedvalues
resid_series = pd.Series(ickmodel.resid)
print("min: ", resid_series.min(), "q1: ", resid_series.quantile(0.25), "median:
↪ ", resid_series.median(), "q3: ", resid_series.quantile(0.75), "max: "
↪ ,resid_series.max())
pd.Series(ickmodel.resid).plot.density();
```

```
min: -0.17585777802014163 q1: -0.06790367526020552 median:
-0.019300475934164396 q3: 0.03725672678702266 max: 1.0106989261669985
```



ok, we're nearly normal, if not a bit spikier around the mean than preferable. Still, so far it looks like there is a linear relationship between vowel duration and metrical position in the song.

Let's see how quantity is shaking out:

```
[ ]: qmodel = smf.ols("seg_duration ~ quantity", data=vowel_df).fit()
      qmodel.summary()
```

```
[ ]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:          seg_duration    R-squared:                0.051
Model:                  OLS            Adj. R-squared:           0.048
Method:                 Least Squares   F-statistic:              16.39
Date:                  Thu, 12 May 2022  Prob (F-statistic):       1.16e-07
Time:                  00:26:19         Log-Likelihood:          461.04
No. Observations:      615             AIC:                   -916.1
Df Residuals:          612             BIC:                   -902.8
Df Model:               2
Covariance Type:       nonrobust
=====
=
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
```

```

-
Intercept          0.2164      0.008      26.428      0.000      0.200
0.232
quantity[T.2]      -0.0107     0.010     -1.026     0.305     -0.031
0.010
quantity[T.3]      -0.0758     0.014     -5.467     0.000     -0.103
-0.049
=====
Omnibus:              508.017   Durbin-Watson:              1.666
Prob(Omnibus):         0.000   Jarque-Bera (JB):          12614.124
Skew:                  3.541   Prob(JB):                  0.00
Kurtosis:              24.026   Cond. No.                  4.12
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

Well, we have significant p values for the first (Q1) and third (Q3) coefficients, but not the second (Q2). This makes sense, since the quantity contrast is indicated by the duration ratios of the syllables. Q3, however, never appears in an unstressed position in a word, so only needs to be contrasted with Q2, while Q1 and Q2 both appear in stressed and unstressed positions and need to be differentiated from each other. The adj r-squared here is a little bit better than the ictus model above, and we do have statistical significance for the model overall ($p < 0.0000$)

```


[ ]: qmodel.fittedvalues
qresid_series = pd.Series(qmodel.resid)
print("min: " ,qresid_series.min(), "q1: " , qresid_series.quantile(0.25),
      ↪ "median: " ,qresid_series.median(), "q3: " , qresid_series.quantile(0.
      ↪ 75), "max: " ,qresid_series.max())
pd.Series(qmodel.resid).plot.density();

```

```

min: -0.20110167790250655 q1: -0.060080153917226674 median:
-0.021591012703996354 q3: 0.0328693256575211 max: 0.9961566923675012

```



```
[ ]: stressmodel = smf.ols("seg_duration ~ stress", data=vowel_df).fit()
stressmodel.summary()
```

```
[ ]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                OLS Regression Results
=====
Dep. Variable:          seg_duration    R-squared:                0.061
Model:                  OLS            Adj. R-squared:           0.059
Method:                 Least Squares   F-statistic:               39.73
Date:                  Thu, 12 May 2022  Prob (F-statistic):       5.58e-10
Time:                  00:26:20         Log-Likelihood:            464.31
No. Observations:      615             AIC:                     -924.6
Df Residuals:          613             BIC:                     -915.8
Df Model:               1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.2286	0.007	34.175	0.000	0.215	0.242
stress	-0.0580	0.009	-6.303	0.000	-0.076	-0.040

```

=====
Omnibus:                476.606    Durbin-Watson:           1.635
Prob(Omnibus):          0.000     Jarque-Bera (JB):        10197.558
Skew:                   3.270     Prob(JB):                 0.00

```

Kurtosis: 21.847 Cond. No. 2.69

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

here we have significant effects for both coefficients. It looks like word-level stress predicts a shorter vowel duration, with a negative slope.

```
[ ]: stress_q_model = smf.ols("seg_duration ~ stress + quantity", data=vowel_df).
    ↪fit()
    stress_q_model.summary()
```

```
[ ]: <class 'statsmodels.iolib.summary.Summary'>
    """
```

OLS Regression Results

```
=====
Dep. Variable:          seg_duration    R-squared:                0.077
Model:                  OLS            Adj. R-squared:           0.073
Method:                 Least Squares   F-statistic:              17.10
Date:                   Thu, 12 May 2022 Prob (F-statistic):       1.12e-10
Time:                   00:43:32        Log-Likelihood:          469.79
No. Observations:       615            AIC:                     -931.6
Df Residuals:           611            BIC:                     -913.9
Df Model:                3
Covariance Type:        nonrobust
=====
```

```
=
              coef    std err          t      P>|t|      [0.025
0.975]
-----
-
Intercept      0.2300      0.009     26.414      0.000      0.213
0.247
stress[T.1]    -0.0432      0.010     -4.200      0.000     -0.063
-0.023
quantity[T.2]  -0.0026      0.010     -0.251      0.802     -0.023
0.018
quantity[T.3]  -0.0462      0.015     -3.007      0.003     -0.076
-0.016
=====
Omnibus:            488.788    Durbin-Watson:           1.653
Prob(Omnibus):      0.000    Jarque-Bera (JB):        11163.529
Skew:               3.370    Prob(JB):                0.00
Kurtosis:           22.754    Cond. No.:               4.99
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

so far, this is the best model we have as far as adj r-squared goes. It is explaining the most variation so far, and it is unlikely with the low p value that the model is complete trash. We didn't lose significance for any of our coefficients from single factor models.

```
[ ]: stress_q_model.fittedvalues
stressresid_series = pd.Series(stress_q_model.resid)
print("min: ", stressresid_series.min(), "q1: ", stressresid_series.quantile(0.
↪25), "median: ", stressresid_series.median(), "q3: ", stressresid_series.
↪quantile(0.75), "max: ", stressresid_series.max())
pd.Series(stress_q_model.resid).plot.density();
```

```
min: -0.19704059472321775 q1: -0.06306568545920294 median:
-0.020152767123178622 q3: 0.03726678408654223 max: 0.974420464328187
```

